

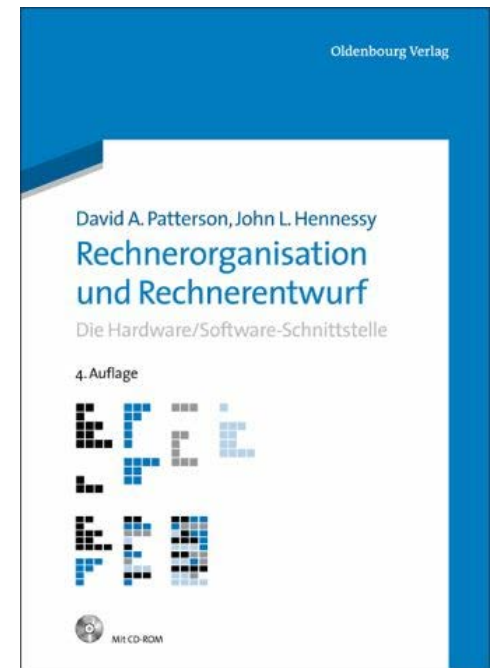
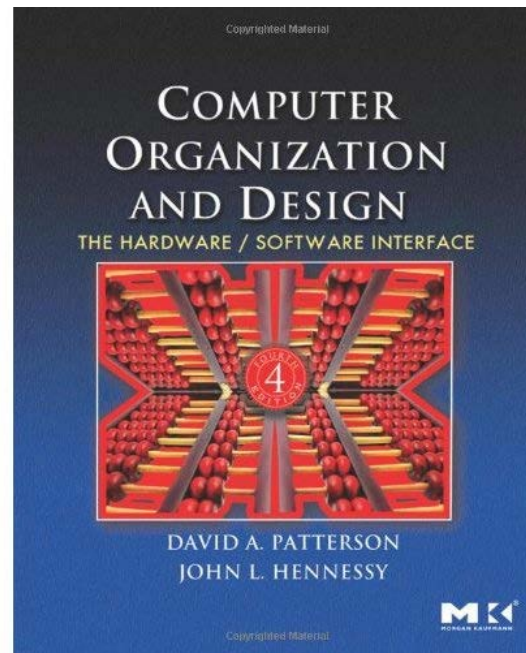


Informatik der Systeme – Chapter 5: Instructions – Language of the Computer

Prof. Dr. Michael Menth

<http://kn.inf.uni-tuebingen.de>

- ▶ This set of slides is an adaptation of Prof. Mary Jane Irwin's lecture notes, <http://www.cse.psu.edu/research/mdl/mji/>
- ▶ Adapted from Patterson & Hennessy: "Computer Organization and Design", 4th Edition, © 2008, MK
- ▶ German translation: Patterson & Hennessy: "Rechnerorganisation und Rechnerentwurf"



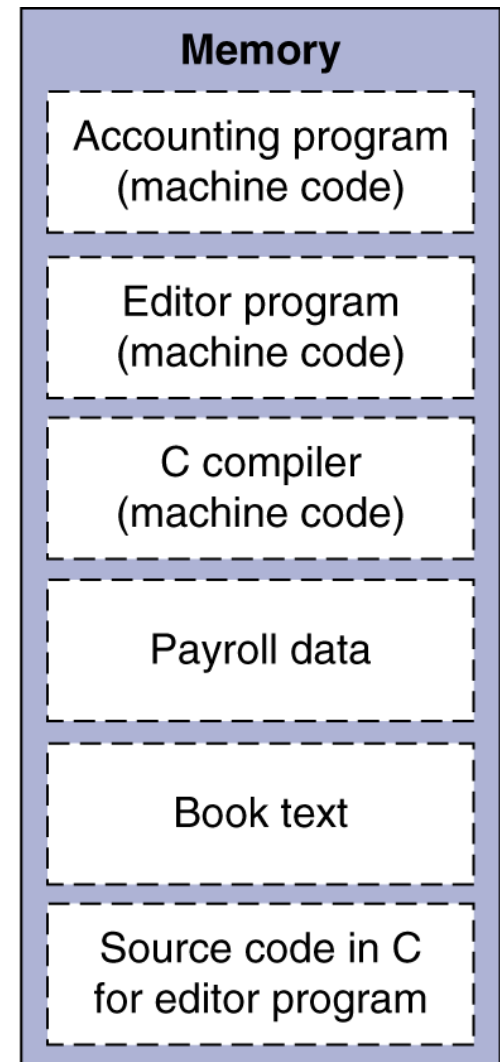


- ▶ Instruction set architecture
- ▶ The MIPS ISA
- ▶ Calling procedures in MIPS
- ▶ Dealing with character sets
- ▶ Translating and starting a program
- ▶ Effect of compiler optimization
- ▶ Other ISAs
- ▶ Concluding remarks



Stored Program Computers (Von Neumann Architecture)

- ▶ Instructions just like data
 - Represented in binary
 - Stored in memory
 - Programs can be shipped as files of binary numbers
- ▶ Programs can operate on programs
 - E.g., compilers, linkers, ...
- ▶ Binary compatibility
 - Allows compiled programs to work on different computers compatible with the same ISA
 - Leads industry to align around a small number of ISAs





- ▶ The repertoire of instructions of a computer
- ▶ Different computers have different instruction sets
 - But with many aspects in common
- ▶ Early computers had very simple instruction sets
 - Simplified implementation
- ▶ Many modern computers also have simple instruction sets, in particular embedded devices



► CISC

- Complex Instruction Set Computer
- Single instructions can execute several low-level operations (such as a load from memory, an arithmetic operation, and a memory store) and/or are capable of multi-step operations or addressing modes within single instructions
- E.g. IBM's System/360 through z/Architecture, PDP-11, VAX, Motorola 68k, and x86 family until 80386

► RISC

- Reduced Instruction Set Computer
- CPU design strategy based on the insight that simplified instructions can provide higher performance if this simplicity enables much faster execution of each instruction.
- E.g. DEC Alpha, AMD 29k, ARC, ARM, Atmel AVR, MIPS, PA-RISC, Power (including PowerPC), and SPARC
- x86: RISC core with CISC emulation
- ARM processors in smart phones and tablet computers such as the iPad



- ▶ Instruction set architecture
- ▶ The MIPS ISA
- ▶ Calling procedures in MIPS
- ▶ Dealing with character sets
- ▶ Translating and starting a program
- ▶ Effect of compiler optimization
- ▶ Other ISAs
- ▶ Concluding remarks



- ▶ Used as the example throughout the lecture
- ▶ **M**icroprocessor without **I**nterlocked **P**ipeline **S**tages
 - RISC-style ISA, developed since 1981 by John Hennessy at Stanford University
 - From 1984 on: MIPS Computer Systems Inc., today MIPS Technologies
- ▶ Large share of embedded core market
 - Applications in consumer electronics, network/storage equipment, cameras, printers, ...
- ▶ Typical of many modern ISAs



► Simplicity favors regularity

- Fixed size instructions
- Small number of instruction formats
- Opcode always the first 6 bits

► Smaller is faster

- Limited instruction set
- Limited number of registers in register file
- Limited number of addressing modes

► Make the common case fast

- Arithmetic operands from the register file (load-store machine)
- Allow instructions to contain immediate operands

► Good design demands good compromises


- Three instruction formats



► Instruction Categories

- Computational
- Load/Store
- Jump and Branch
- Floating Point
 - Coprocessor
- Memory Management
- Special

Registers (32 bits wide)



3 Instruction Formats (machine code representation): all 32 bits wide

op	rs	rt	rd	sa	funct	R format
op	rs	rt	immediate			I format
op	jump target					J format

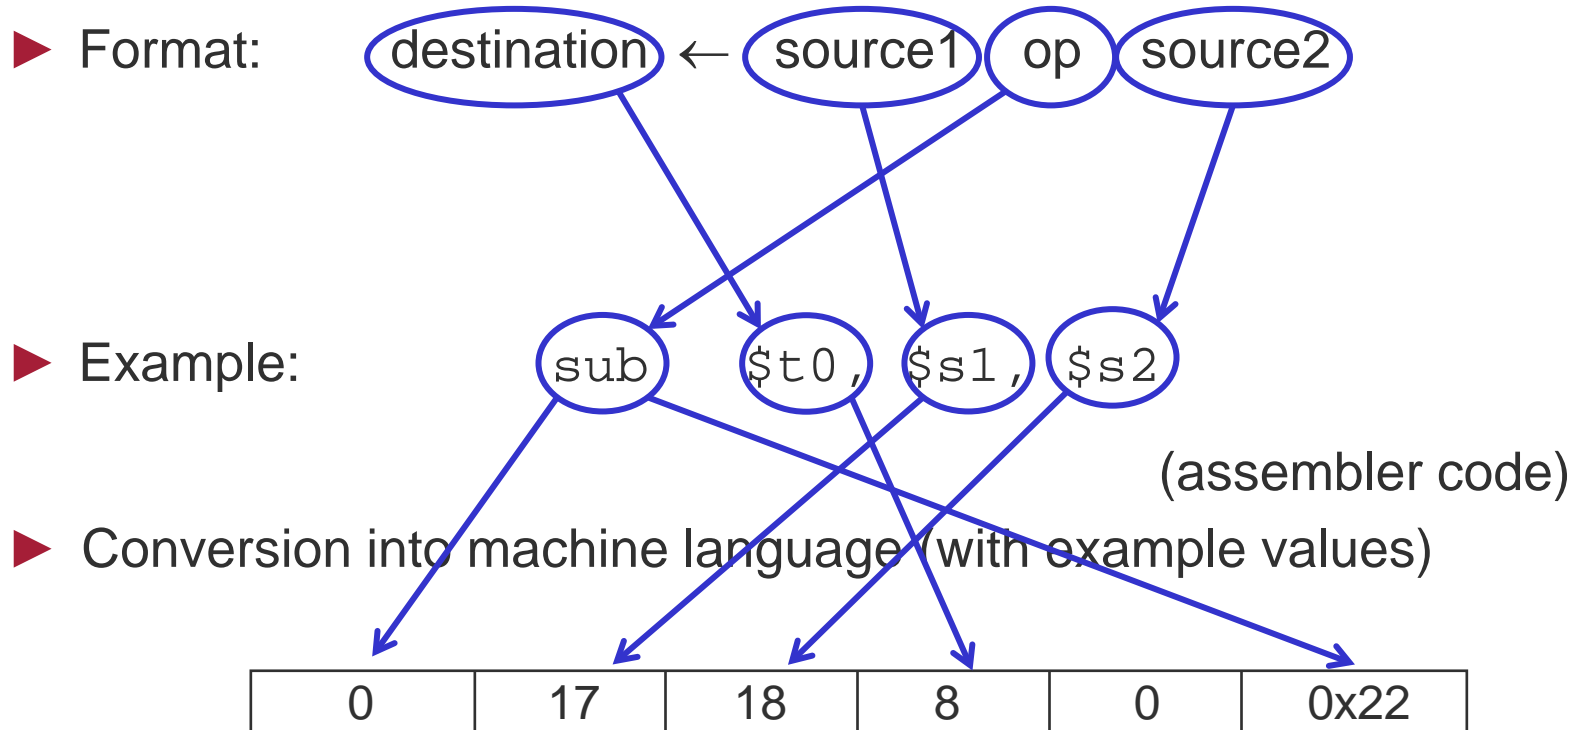


Aside: MIPS Register Convention

Name	Register Number	Usage	Preserve on call?
\$zero	0	constant 0 (hardware)	n.a.
\$at	1	reserved for assembler	n.a.
\$v0 - \$v1	2-3	returned values	no
\$a0 - \$a3	4-7	arguments	yes
\$t0 - \$t7	8-15	temporaries	no
\$s0 - \$s7	16-23	saved values	yes
\$t8 - \$t9	24-25	temporaries	no
\$gp	28	global pointer	yes
\$sp	29	stack pointer	yes
\$fp	30	frame pointer	yes
\$ra	31	return addr (hardware)	yes



- ▶ Each arithmetic instruction in MIPS assembly language performs one operation
- ▶ Each specifies exactly three operands that are all contained in the datapath's register file (\$t0, \$s1, \$s2)





► MIPS instruction formats define special instruction fields

► **R**-Format

- For instructions with up to 3 **R**egisters

op	rs	rt	rd	shamt	funct
----	----	----	----	-------	-------

- op 6-bits opcode that specifies the operation
- rs 5-bits register file address of the first source operand
- rt 5-bits register file address of the second source operand
- rd 5-bits register file address of the result's destination
- shamt 5-bits shift amount (for shift instructions)
- funct 6-bits function code augmenting the opcode

► **I**-Format (later)

- For instructions with **I**mmEDIATE operands (explicit numbers)

► **J**-Format (later)

- For **J**ump instructions



op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

add \$t0, \$s1, \$s2

special	\$s1	\$s2	\$t0	0	add
0	17	18	8	0	32
000000	10001	10010	01000	00000	100000

$$00000010001100100100000000100000_2 = 02324020_{16}$$

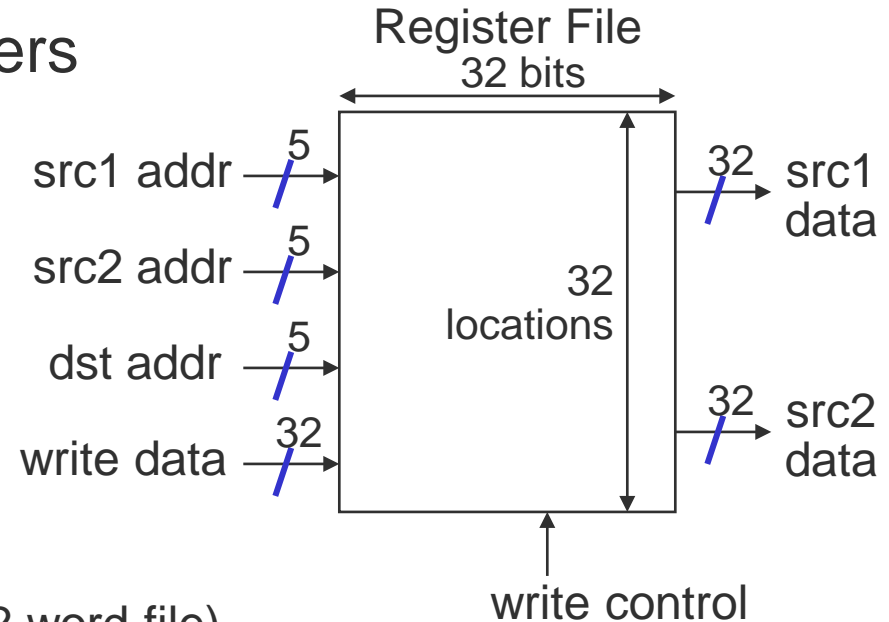


► Holds thirty-two 32-bit registers

- Two read ports and
- One write port


► Registers are

- Faster than main memory
 - But register files with more locations are slower (e.g., a 64 word file could be as much as 50% slower than a 32 word file)
 - Read/write port increase impacts speed quadratically
- Easier for a compiler to use
 - E.g., $(A \cdot B) - (C \cdot D) - (E \cdot F)$ can do multiplies in any order vs. stack
- Can hold variables so that
 - Code density improves (since register are named with fewer bits than a memory location)





- ▶ Used for composite data (variables, arrays, structures, dynamic data, ...)
- ▶ To apply arithmetic operations
 - Load values from memory into registers
 - Store result from register to memory
- ▶ Most architectures address individual **bytes** in memory
- ▶ **Alignment restriction** – the memory address of a **word** must be on natural word boundaries (a multiple of 4 in MIPS-32)
- ▶ **Big Endian**: leftmost byte is word address
 - Examples: IBM 360/370, Motorola 68k, **MIPS**, Sparc, HP PA
- ▶ **Little Endian**: rightmost byte is word address
 - Examples: Intel 80x86, DEC Vax, DEC Alpha (Windows NT)



„Gulliver’s Travels“, Jonathan Swift



MIPS Memory Access Instructions

- ▶ MIPS has two basic **data transfer instructions** for accessing memory

```
lw    $t0, 4($s3) #load word from memory
```

```
sw    $t0, 8($s3) #store word to memory
```

- ▶ **I-Format** – for immediate arithmetic and load/store instructions

- rt: destination or source register number
- Memory address: sum of offset and base address in rs




- ▶ The data is loaded into (lw) or stored from (sw) a register in the register file (\$t0) – a 5 bit address
- ▶ The memory address – a 32 bit address – is formed by adding the contents of the **base address register** (\$s3) to the 16 bit **offset** value
 - A 16-bit field meaning access is limited to memory locations within a region of $\pm 2^{13}$ or 8,192 words ($\pm 2^{15}$ or 32,768 bytes) of the address in the base register




Example: Load Instruction

► Load/store instructions use I-Format



$$24_{10} + \$s3 =$$

$$\begin{array}{r}
 \dots 0001\ 1000 \\
 + \dots 1001\ 0100 \\
 \hline
 \dots 1010\ 1100 = \\
 \quad 0x120040ac
 \end{array}$$



Word addresses are multiples of 4!



► MIPS provides special instructions to move **bytes**

```
lb    $t0, 1($s3)    #load byte from memory
sb    $t0, 6($s3)    #store byte to  memory
```

0x28	19	8	16 bit offset
------	----	---	---------------

► What 8 bits get loaded and stored?

- Load byte places the byte from memory in the rightmost 8 bits of the destination register
 - What happens to the other bits in the register?
- Store byte takes the byte from the rightmost 8 bits of a register and writes it to a byte in memory
 - What happens to the other bits in the memory word?



► Small constants are used often in typical code

► Possible approaches?

- Put “typical constants” in memory and load them
- Create hard-wired registers (like \$zero) for constants like 1
- Have special instructions that contain constants !

`addi $s0, $s0, 4 # $s0 = $s0 + 4`

`slti $t0, $s2, 15 # $t0 = 1 if $s2 < 15`

- Immediate instructions use I-Format



► The constant is kept inside the instruction itself!

- Immediate format limits values to the range from -2^{15} to $+2^{15}-1$



Aside: How About Larger Constants?

- ▶ We'd also like to be able to load a 32 bit constant into a register, for this we must use two instructions

- ▶ A new "load upper immediate" instruction

```
lui $t0, 1010101010101010
```

16	0	8	1010101010101010 ₂
----	---	---	-------------------------------

- ▶ Then must get the lower order bits right, use

```
ori $t0, $t0, 1010101010101010
```

1010101010101010	0000000000000000
------------------	------------------

0000000000000000	1010101010101010
------------------	------------------

1010101010101010	1010101010101010
------------------	------------------



► Instructions for bitwise manipulation

- Useful for extracting and inserting groups of bits (e.g. bytes or 8-bit characters) in a 32-bit word
- Shift left logical (sll): shift left and fill with 0 bits
- Shift right logical (srl): shift right and fill with 0 bits
- Useful for fast multiplication of unsigned numbers by 2^i

Operation	C	Java	MIPS
Shift left	<<	<<	sll
Shift right	>>	>>	srl
Bitwise AND	&	&	and, andi
Bitwise OR			or, ori
Bitwise NOT	~	~	nor



- Shifts move all the bits in a word left or right

`sll $t2, $s0, 8` `#$t2 = $s0 << 8 bits`

`srl $t2, $s0, 8` `#$t2 = $s0 >> 8 bits`

- Shift operations use R-Format



- Shift amount (shamt) = 8
- Notice that a 5-bit shamt field is enough to shift a 32-bit value 2^5 – 1 or 31 bit positions



- Useful to mask bits in a word
 - Select some bits, clear others to 0

and \$t0, \$t1, \$t2

\$t2	0000	0000	0000	0000	00	00	11	01	11	00	0000
\$t1	0000	0000	0000	0000	00	11	11	00	0000	0000	
\$t0	0000	0000	0000	0000	00	00	11	00	0000	0000	



- Useful to include bits in a word
 - Set some bits to 1, leave others unchanged

or \$t0, \$t1, \$t2

\$t2	0000	0000	0000	0000	00	00	11	01	11	00	0000
\$t1	0000	0000	0000	0000	00	11	11	00	0000	0000	
\$t0	0000	0000	0000	0000	00	11	11	01	11	00	0000



- ▶ Useful to invert bits in a word
 - Change 0 to 1, and 1 to 0
- ▶ MIPS has NOR 3-operand instruction
 - $a \text{ NOR } b == \text{NOT} (a \text{ OR } b)$

Register 0: always
read as zero

nor \$t0, \$t1, \$zero

\$zero	0000 0000 0000 0000 0000 0000 0000 0000
\$t1	0000 0000 0000 0000 0011 1100 0000 0000
\$t0	1111 1111 1111 1111 1100 0011 1111 1111



- There are a number of **bit-wise** logical operations in the MIPS ISA

`and $t0, $t1, $t2` $\# \$t0 = \$t1 \ \& \ \$t2$

`or $t0, $t1, $t2` $\# \$t0 = \$t1 \ | \ \$t2$

`nor $t0, $t1, $t2` $\# \$t0 = \text{not}(\$t1 \ | \ \$t2)$

- Logical operations with two source operands use **R-Format**

0	9	10	8	0	0x24
---	---	----	---	---	------

`andi $t0, $t1, 0xFF00` $\# \$t0 = \$t1 \ \& \ \text{ff00}$

`ori $t0, $t1, 0xFF00` $\# \$t0 = \$t1 \ | \ \text{ff00}$

- Logical operations with immediate operand use **I-Format**

0x0D	9	8	0xFF00
------	---	---	--------



► MIPS conditional branch instructions

```
bne $s0, $s1, Lbl      #go to Lbl if $s0≠$s1
beq $s0, $s1, Lbl      #go to Lbl if $s0=$s1
```

► Example

```
        if (i==j) h = i + j;
        bne $s0, $s1, Lbl1
        add $s3, $s0, $s1
Lbl1:    ...
```

► Assembler takes care of fixing the label

- Lbl1 in the instruction is approximately the offset to the current instruction address (more next slide)
- “Lbl1:” marks the corresponding instruction in code

► Conditional branch uses I-Format

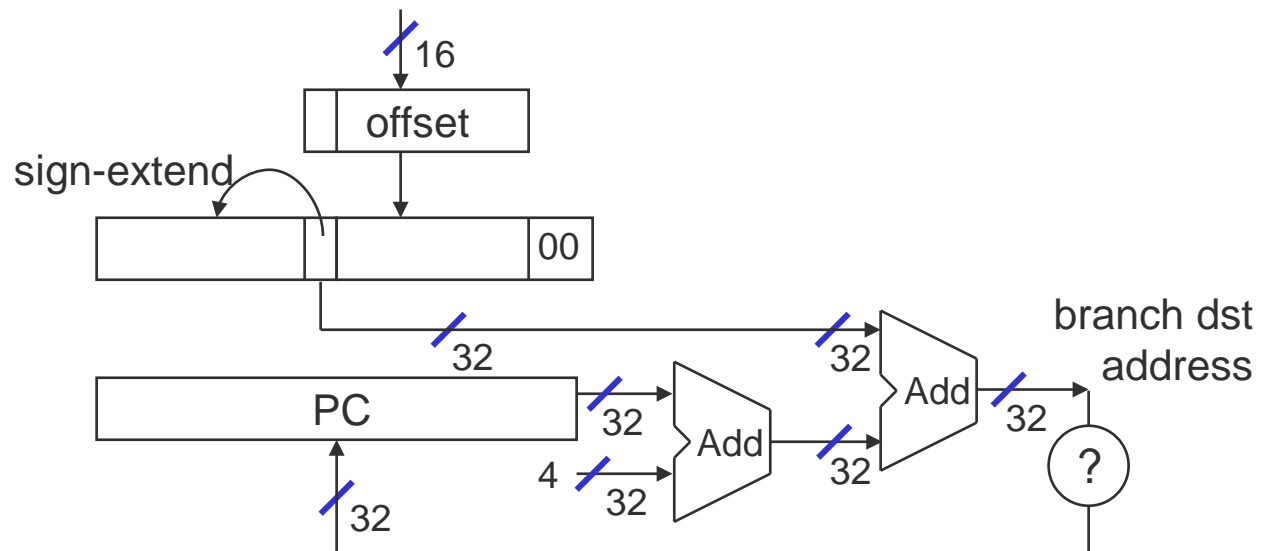
0x05	16	17	16 bit offset
------	----	----	---------------



Specifying Branch Destinations


- ▶ Example: `bne $s0, $s1, Lbl1`
- ▶ Means: next instruction address is $(PC+4)+4*(\text{value of Lbl1})$
 - PC register automatically incremented ($PC+4$) during fetch cycle
- ▶ Limits the branch distance to -2^{15} to $+2^{15}-1$ (word) instructions from the (instruction after the) branch instruction
 - But most branches are local anyway

from the low order 16 bits of the branch instruction





- ▶ A basic block is a sequence of instructions with
 - No embedded branches (except at end)
 - No branch targets (except at beginning)
- ▶ A compiler identifies basic blocks for optimization
- ▶ An advanced processor can accelerate execution of basic blocks





- ▶ We have `beq`, `bne`, but what about other kinds of branches (e.g., branch-if-less-than)?
- ▶ Set on less than instruction

```
slt $t0, $s0, $s1      # if $s0 < $s1      then
                        # $t0 = 1            else
                        # $t0 = 0
```

- Instruction format (R format):

0	16	17	8		0x24
---	----	----	---	--	------

- ▶ Immediate operand available

```
slti $t0, $s0, 25      # if $s0 < 25 then $t0=1 ...
```



Signed vs. Unsigned Comparison

- Unsigned versions of `slt` and `slti` (u stands for unsigned)

```
sltu $t0, $s0, $s1    # if $s0 < $s1 then $t0=1 ...
```

```
sltiu $t0, $s0, 25    # if $s0 < 25 then $t0=1 ...
```

- Example

- `$s0 = 1111 1111 1111 1111 1111 1111 1111 1111`
- `$s1 = 0000 0000 0000 0000 0000 0000 0000 0001`
- `slt $t0, $s0, $s1 # signed`
 - $-1 < +1 \Rightarrow \$t0 = 1$
- `sltu $t0, $s0, $s1 # unsigned`
 - $+4,294,967,295 > +1 \Rightarrow \$t0 = 0$



Aside: More Branch Instructions

- ▶ Why not `blt`, `bge`, etc? Hardware for $<$, \geq , ... slower than $=$, \neq
 - Combining with branch involves more work per instruction, requiring a slower clock – all instructions penalized!
- ▶ Can use `slt`, `beq`, `bne`, and the fixed value of 0 in register `$zero` to create other conditions
 - Less than


```
blt $s1, $s2, Label
    slt $at, $s1, $s2      # $at set to 1 if
    bne $at, $zero, Label  # $s1 < $s2
```
 - Less than or equal to


```
ble $s1, $s2, Label
```
 - Greater than


```
bgt $s1, $s2, Label
```
 - Greater than or equal to


```
bge $s1, $s2, Label
```
- ▶ Such branches are included in the instruction set as pseudo instructions
 - Recognized (and expanded) by the assembler
 - Therefore, the assembler needs a reserved register (`$at`)



- ▶ Treating signed numbers as if they were unsigned gives a low-cost way of checking if $0 \leq x < y$ (index out of bounds for arrays)

- $y \rightarrow \$t2, x \rightarrow \$s1$

```
sltu $t0, $s1, $t2    # $t0 = 1 if
                      # $s1 < $t2 (max)
                      # and $s1 >= 0 (min)
beq $t0, $zero, IOOB  # go to IOOB if
                      # $t0 = 0
```

- ▶ The key is that negative integers in two's complement look like large numbers in unsigned notation. Thus, an unsigned comparison of $x < y$ also checks if x is negative as well as if x is less than y .




- MIPS also has an unconditional branch instruction or **jump** instruction:

`j label` # go to label

- Jump instruction uses the **J-Format**



from the low order 26 bits of the jump instruction





► C code:

```
while (save[i] == k) i += 1;
```


- i in \$s3, k in \$s5, address of save in \$s6

► Compiled MIPS code:

```

Loop: slt $t1, $s3, 2
      add $t1, $t1, $s6
      lw  $t0, 0($t1)
      bne $t0, $s5, Exit
      addi $s3, $s3, 1
      j   Loop
Exit: ...

```





Aside: Conditional Branching Far Away

- ▶ What if a **conditional** branch destination is further away than can be captured in 16 bits?
- ▶ The assembler comes to the rescue – it inserts an unconditional jump to the branch target and inverts the condition

▶ Example: `beq $s0, $s1, L1`

▶ Becomes:

```

    bne    $s0, $s1, L2
    j      L1
L2:    ...
  
```



► MIPS **procedure call** instruction

`jal ProcedureAddress #jump and link`

- Saves PC+4 in register `$ra` to have a link to the next instruction for the procedure return
- Uses the **J-Format**

0x03	26 bit address
------	----------------

► Then can do procedure **return** with a

`jr $ra #return`

- Can also be used for computed jumps
 - e.g., for case/switch statements
- Uses the **R-Format**

0	31				0x08
---	----	--	--	--	------



► Must support the arithmetic/logic operations of the ISA

`add, addi, addiu, addu`

`sub, subu`

`mult, multu, div, divu`

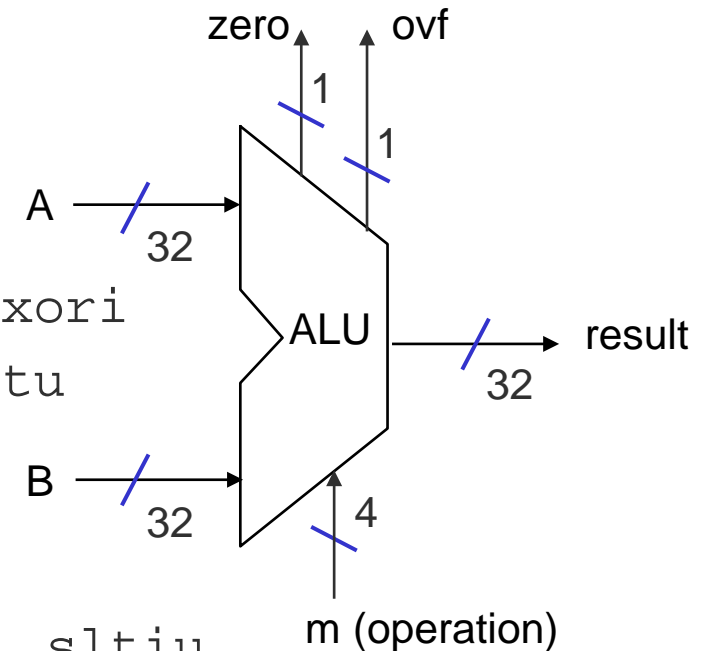
`sqrt`

`and, andi, nor, or, ori, xor, xori`

`beq, bne, slt, slti, sltiu, sltu`

► With special handling for

- Sign extend – `addi, addiu, slti, sltiu`
- Zero extend – `andi, ori, xori`
- Overflow detection – `add, addi, sub`





- ▶ Overflow occurs when the result of an operation cannot be represented in 32-bits, i.e., when the sign bit contains a value bit of the result and not the proper sign bit
- ▶ When adding operands with different signs or when subtracting operands with the same sign, overflow can never occur

Operation	Operand A	Operand B	Result indicating overflow
$A + B$	≥ 0	≥ 0	< 0
$A + B$	< 0	< 0	≥ 0
$A - B$	≥ 0	< 0	< 0
$A - B$	< 0	≥ 0	≥ 0

- ▶ MIPS signals overflow with an exception (aka interrupt) – an unscheduled procedure call where the EPC contains the address of the instruction that caused the exception



▶ Register operands

- Direct access
- Used by arithmetic instructions
- Only 32 registers to enable fast access in implementations

▶ Memory operands

- Only accessible via load and store instructions
 - ⇒ Compilers must use registers for variables as much as possible
- Only spill to memory for less frequently used variables
- Register optimization is important!

▶ Immediate operands

- Explicit numbers in instructions

▶ MIPS register 0 (\$zero) is the constant 0

- Cannot be overwritten
- Useful for common operations

`add $t2, $s1, $zero`



Summary of Addressing Modes in MIPS

► Immediate addressing

- `addi $s0 $s1 1`

► Register addressing

- `add $s0 $s1 $s2`

► Base addressing

- `sw $s0 $s3(4)`


► PC-relative addressing

- `bne $s0 $s1 loop`


► Pseudo-direct addressing

- `j loop`

1. Immediate addressing





2. Register addressing





Registers
Register



3. Base addressing



4. PC-relative addressing



5. Pseudodirect addressing





- ▶ Instruction set architecture
- ▶ The MIPS ISA
- ▶ Calling procedures in MIPS
- ▶ Dealing with character sets
- ▶ Translating and starting a program
- ▶ Effect of compiler optimization
- ▶ Other ISAs
- ▶ Concluding remarks



1. Main routine (**caller**) places parameters in a place where the procedure (**callee**) can access them
 - `$a0 – $a3`: four **argument registers**
2. Caller transfers control to the callee
3. Callee acquires the storage resources needed
4. Callee performs the desired task
5. Callee places the result value in a place where the caller can access it
 - `$v0 – $v1`: two **value registers** for result values
6. Callee returns control to the caller
 - `$ra`: one **return address register** to return to the point of origin



Spilling Registers into Memory

- What if the callee needs to use more registers than allocated to argument and return values?

- Callee uses a **stack** – a last-in-first-out queue


- Stack pointer $\$sp$ ($\$29$)

- One of the general registers
 - Used to address the stack
 - Stack “grows” from high address to low address
 - Add data onto the stack – push data on stack at new $\$sp$

$$\$sp = \$sp - 4$$

- Remove data from the stack – pop data from stack at $\$sp$

$$\$sp = \$sp + 4$$





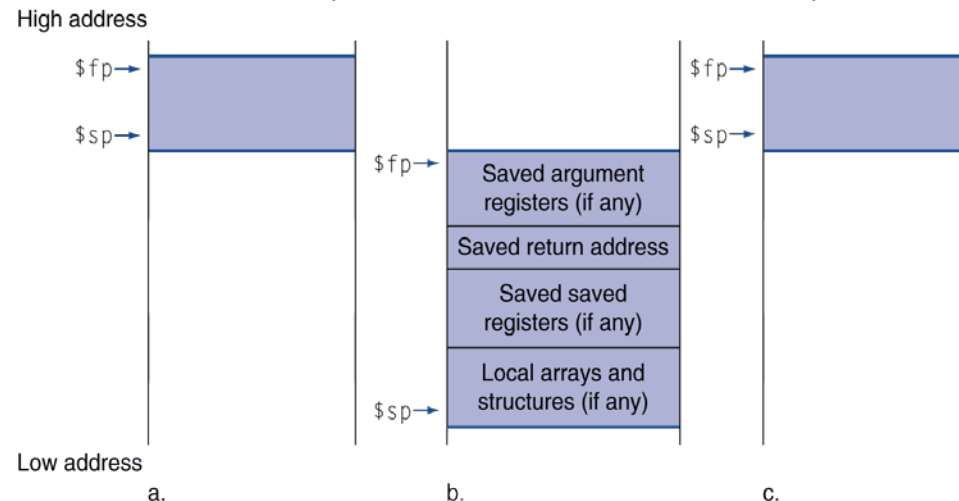

Aside: Allocating Space on the Stack

► Procedure frame

- The segment of the stack containing a procedure's saved registers and local variables
- Aka **activation record**

► The frame pointer (\$fp)

- Points to the first word of the frame of a procedure
- Provides a stable “base” register for the procedure
 - \$fp is initialized using \$sp on a call and
 - \$sp is restored using \$fp on a return
- Right figure: before, during, and after procedure call





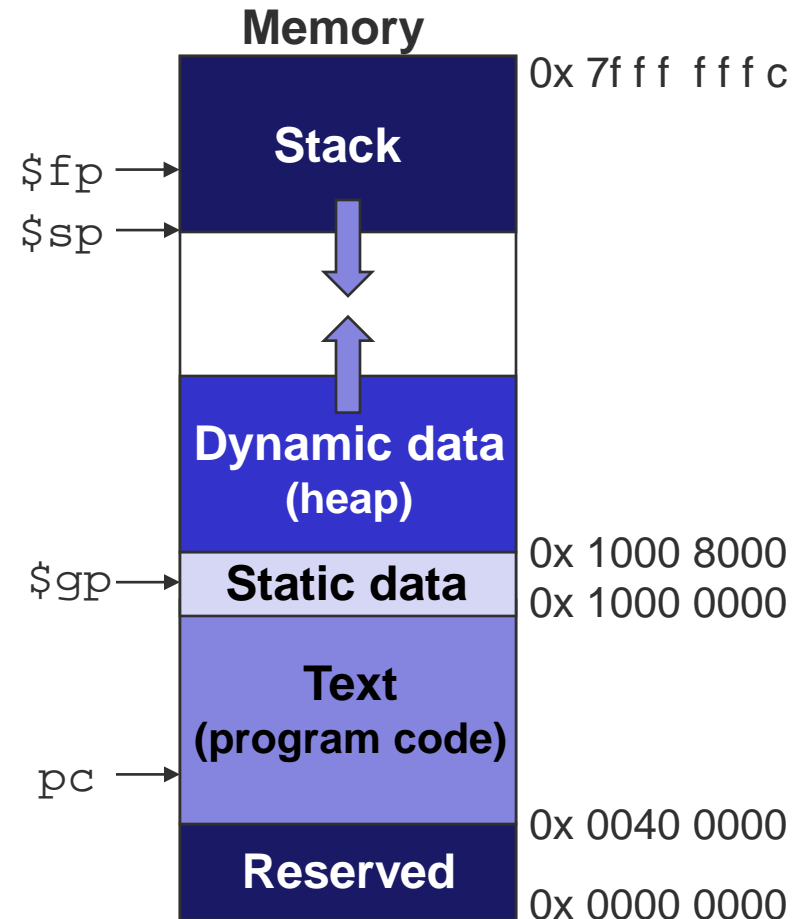
Aside: Allocating Space on the Heap

► Static data segment

- For constants and other static variables (e.g., arrays)
- Global pointer `$gp` used to address static data with positive and negative 16-bit offset (`lw`, `sw`)

► Dynamic data segment

- For structures that grow and shrink (e.g., linked lists)
- Aka **heap**
- Dealing with dynamic data in C
 - Allocate space on the heap with `malloc()` (memory allocation)
 - Free it with `free()`





► Procedures

- Must save register values $\$sx$ on the stack if it makes use of these registers
- Must restore them before return
- Can overwrite arguments $\$ax$ and temporaries $\$tx$

► Leaf procedure

- Does not call other procedures

► Non-leaf procedure

- Calls other procedures (nested calls)
- For nested call, caller needs to save on the stack
 - Its return address
 - Any arguments $\$ax$ and temporaries $\$tx$ needed after the call
- Registers restored from the stack after the call



► C code

```
int leaf_example (int g, h, i, j)
{ int f;
  f = (g + h) - (i + j);
  return f;}
```

► MIPS code

- Arguments g, ..., j in \$a0, ..., \$a3
- f in \$s0 (hence, need to save \$s0 on stack)
- Result in \$v0

```
leaf_example:
    addi $sp, $sp, -4      # Save $s0 on stack
    sw   $s0, 0($sp)
    add  $t0, $a0, $a1     # Procedure body
    add  $t1, $a2, $a3
    sub  $s0, $t0, $t1
    add  $v0, $s0, $zero   # Result
    lw   $s0, 0($sp)      # Restore $s0
    addi $sp, $sp, 4
    jr   $ra              # Return
```



Non-Leaf Procedure Example

► C code

```
int fact (int n){
    if (n < 1) return 1;
    else return n * fact(n - 1);}
```

Recursive procedure:
special case of a non-leaf
procedure that calls itself again.

► MIPS code: argument n in \$a0, result in \$v0

fact:

```
        addi $sp, $sp, -8      # adjust stack for 2 items
        sw   $ra, 4($sp)      # save return address
        sw   $a0, 0($sp)      # save argument
        slti $t0, $a0, 1      # test for n < 1
        beq  $t0, $zero, L1
        addi $v0, $zero, 1     # if so, result is 1
        addi $sp, $sp, 8      # pop 2 items from stack
        jr   $ra              # and return
L1:     addi $a0, $a0, -1      # else decrement n
        jal  fact              # recursive call
        lw   $a0, 0($sp)      # restore original n
        lw   $ra, 4($sp)      # and return address
        addi $sp, $sp, 8      # pop 2 items from stack
        mul  $v0, $a0, $v0     # multiply to get result
        jr   $ra              # and return
```




Runtime Example (Requires Animation)

Function calls	Stack occupation	Stack Pointer	Calculation of return values
		$\leftarrow \$sp$	$\$v0=6$
fact(3)	$\$ra$		$\$v0=\$v0*\$a0=2*3=6$
	$\$a0=3$	$\leftarrow \$sp$	$\$v0=2$
fact(2)	$\$ra$		$\$v0=\$v0*\$a0=1*2=2$
	$\$a0=2$	$\leftarrow \$sp$	$\$v0=1$
fact(1)	$\$ra$		$\$v0=\$v0*\$a0=1*1=1$
	$\$a0=1$	$\leftarrow \$sp$	$\$v0=1$
fact(0)	$\$ra$		$\$v0=1$
	$\$a0=0$	$\leftarrow \$sp$	



- Illustrates use of assembly instructions for a C **bubble (exchange) sort** function
- Swap procedure (leaf)

```
void swap(int v[], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```





Example: “swap” in MIPS Assembler

► v in \$a0, k in \$a1, temp in \$t0

swap:	sll \$t1, \$a1, 2	# \$t1 = k * 4
	add \$t1, \$a0, \$t1	# \$t1 = v+(k*4)
		# (address of v[k])
	lw \$t0, 0(\$t1)	# \$t0 (temp) = v[k]
	lw \$t2, 4(\$t1)	# \$t2 = v[k+1]
	sw \$t2, 0(\$t1)	# v[k] = \$t2 (v[k+1])
	sw \$t0, 4(\$t1)	# v[k+1] = \$t0 (temp)
	jr \$ra	# return to calling routine



► Non-leaf (calls swap)

- `v`: array to be sorted
- `n`: size of array to be sorted

```
void sort (int v[], int n)
{
    int i, j;
    for (i = 1*; i < n; i += 1) {
        for (j = i - 1;
             j >= 0 && v[j] > v[j + 1];
             j -= 1) {
            swap(v, j);
        }
    }
}
```

**Book says `i=0`, but `i=1` should be more efficient.*



- ▶ $n=4$ array elements; array positions: 0, 1, 2, 3
- ▶ $v=[4,3,2,1]$ (example)

- ▶ $i=1$
 - $j=0$: $v=[3,4,2,1]$

- ▶ $i=2$
 - $j=1$: $v=[3,2,4,1]$
 - $j=0$: $v=[2,3,4,1]$

- ▶ $i=3$
 - $j=2$: $v=[2,3,1,4]$
 - $j=1$: $v=[2,1,3,4]$
 - $j=0$: $v=[1,2,3,4]$



Example: Procedure Body of “sort” in MIPS

► v in \$a0, n in \$a1, i in \$s0, j in \$s1

	move \$s2, \$a0	# save \$a0 into \$s2	Move params
	move \$s3, \$a1	# save \$a1 into \$s3	
	addi \$s0, \$zero, 1	# i = 1	Outer loop
for1tst:	slt \$t0, \$s0, \$s3	# \$t0 = 1 if \$s0 < \$s3 (i < n)	
	beq \$t0, \$zero, exit1	# go to exit1 if \$s0 ≥ \$s3 (i ≥ n)	
	addi \$s1, \$s0, -1	# j = i - 1	
for2tst:	slti \$t0, \$s1, 0	# \$t0 = 1 if \$s1 < 0 (j < 0)	
	bne \$t0, \$zero, exit2	# go to exit2 if \$s1 < 0 (j < 0)	
	sll \$t1, \$s1, 2	# \$t1 = j * 4	Inner loop
	add \$t2, \$s2, \$t1	# \$t2 = v + (j * 4)	
	lw \$t3, 0(\$t2)	# \$t3 = v[j]	
	lw \$t4, 4(\$t2)	# \$t4 = v[j + 1]	
	slt \$t0, \$t4, \$t3	# \$t0 = 1 if \$t4 < \$t3 (v[j+1] < v[j])	
	beq \$t0, \$zero, exit2	# go to exit2 if \$t4 ≥ \$t3	
	move \$a0, \$s2	# 1st param of swap is v (old \$a0)	Pass params & call
	move \$a1, \$s1	# 2nd param of swap is j	
	jal swap	# call swap procedure	
	addi \$s1, \$s1, -1	# j -= 1	Inner loop
	j for2tst	# jump to test of inner loop	
exit2:	addi \$s0, \$s0, 1	# i += 1	Outer loop
	j for1tst	# jump to test of outer loop	



Example: Full Procedure of “sort” in MIPS

sort:	addi \$sp, \$sp, -20	# make room on stack for 5 registers
	sw \$ra, 16(\$sp)	# save \$ra on stack
	sw \$s3, 12(\$sp)	# save \$s3 on stack
	sw \$s2, 8(\$sp)	# save \$s2 on stack
	sw \$s1, 4(\$sp)	# save \$s1 on stack
	sw \$s0, 0(\$sp)	# save \$s0 on stack
	...	# procedure body
	...	
exit1:	lw \$s0, 0(\$sp)	# restore \$s0 from stack
	lw \$s1, 4(\$sp)	# restore \$s1 from stack
	lw \$s2, 8(\$sp)	# restore \$s2 from stack
	lw \$s3, 12(\$sp)	# restore \$s3 from stack
	lw \$ra, 16(\$sp)	# restore \$ra from stack
	addi \$sp, \$sp, 20	# restore stack pointer
	jr \$ra	# return to calling routine



- ▶ Instruction set architecture
- ▶ The MIPS ISA
- ▶ Calling procedures in MIPS
- ▶ Dealing with character sets
- ▶ Translating and starting a program
- ▶ Effect of compiler optimization
- ▶ Other ISAs
- ▶ Concluding remarks



- ▶ Byte-encoded character sets
 - ASCII: 128 characters, 95 graphic, 33 control
 - Latin-1 (aka ISO 8859-1): 256 characters, ASCII, +96 more graphic characters
- ▶ Unicode (16 bit)
 - Most of the world's alphabets, plus symbols
 - 17x16-bit character sets (1114112 codepoints)
 - Java uses 16 bit Unicode (UTF-16)
 - UTF: Unicode Transformation Format, also UCS Transformation Format (UCS = Universal Character Set)
 - Variable-length encoding: special characters require several bytes (1-4 bytes)
 - Represent all Unicode characters
 - Several variants: UTF-8, UTF-16, UTF-32, UTF-EBCDIC, ...
- ▶ EBCDIC (Extended Binary Coded Decimals Interchange Code)
 - Mostly used on mainframes



- ▶ Array of characters, first position(s) used to indicate array length
 - Used in Java
 - First word indicates length
 - 16-bit Unicode characters used
- ▶ Structure, consisting of a length field and an array of that length
- ▶ Array of characters, end of string determined by value 0
 - Example: ASCII value 0 means symbol „null“ while value 48 means character „0“ \Rightarrow character „0“ can be used in strings
 - Principle used in C
 - Character representation depends on machine and operating system
 - ASCII, OEM, ANSI, ISO-Latin-1, Unicode, EBCDIC,...
 - Most of them use 1 byte for character representation



- ▶ Could use bitwise operations

- ▶ MIPS byte/halfword load/store
 - String processing is a common case
 - Sign extend to 32 bits in rt
 - `lb rt, offset(rs)`
 - `lh rt, offset(rs)`
 - Zero extend to 32 bits in rt
 - `lbu rt, offset(rs)`
 - `lhu rt, offset(rs)`
 - Store just rightmost byte/halfword
 - `sb rt, offset(rs)`
 - `sh rt, offset(rs)`



- ▶ C code (naïve) for copying null-terminated strings
 - Copies string `y[]` into string `x[]`
 - `do ... while(cond)` in C language corresponds to
 - `repeat ... until(not(cond))` in other programming languages

```
void strcpy (char x[], char y[]){
    int i;
    i=-1;
    do{
        i=i+1;
        x[i]=y[i];
    }while (y[i]!='\0')
}
```



String Copy Example in MIPS

- ▶ Addresses of `x`, `y` in `$a0`, `$a1`, value of `i` in `$s0`
- ▶ Character `'\0'` has value 0 in ASCII

`strcpy:`

```

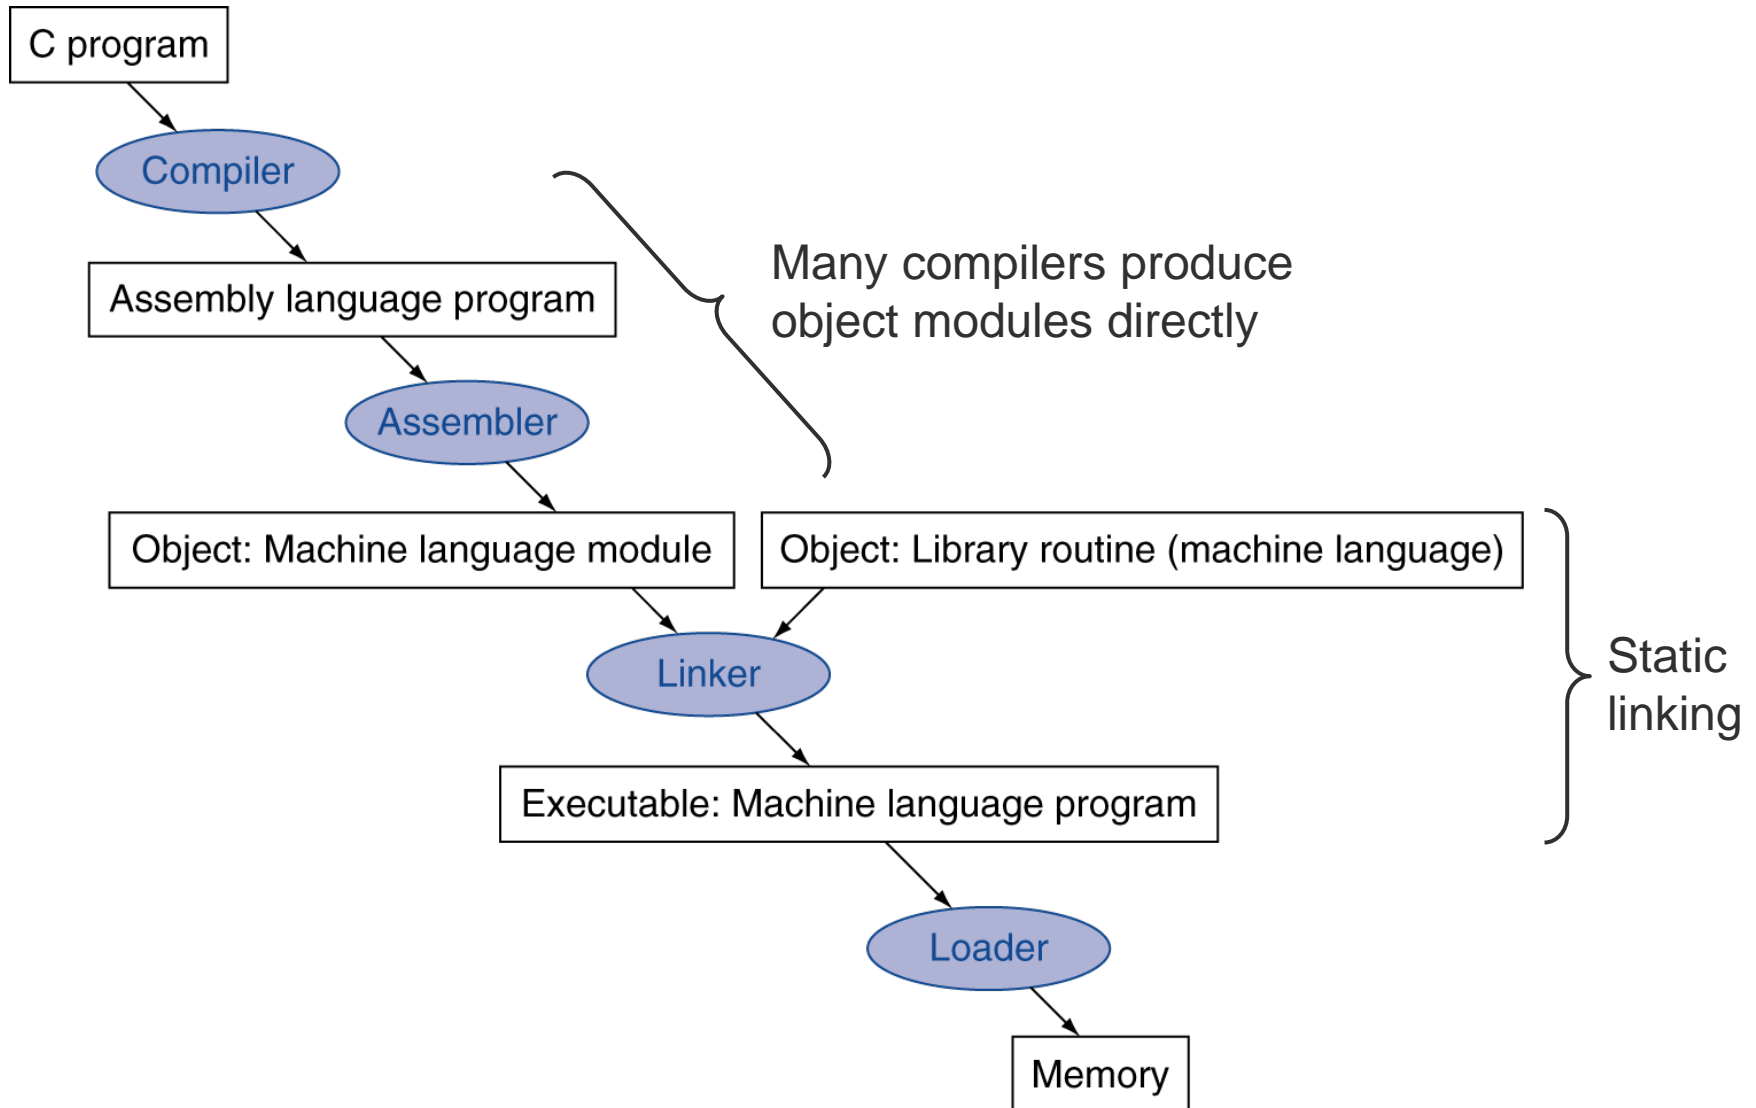
    addi $sp, $sp, -4      # adjust stack for 1 item
    sw    $s0, 0($sp)      # save $s0
    addi  $s0, $zero, -1   # i = -1

# start of do ... while statement
do: addi $s0, $s0, 1       # i = i + 1 (bytes, not words!)
    add   $t1, $s0, $a1    # addr of y[i] in $t1
    lbu   $t2, 0($t1)      # $t2 = y[i]
    add   $t3, $s0, $a0    # addr of x[i] in $t3
    sb    $t2, 0($t3)      # x[i] = y[i]
    bne   $t2, $zero, do   # continue loop if y[i] != 0
# end of do ... while statement
    lw    $s0, 0($sp)      # restore saved $s0
    addi  $sp, $sp, 4      # pop 1 item from stack
    jr    $ra              # and return

```



- ▶ Instruction set architecture
- ▶ The MIPS ISA
- ▶ Calling procedures in MIPS
- ▶ Dealing with character sets
- ▶ Translating and starting a program
- ▶ Effect of compiler optimization
- ▶ Other ISAs
- ▶ Concluding remarks





- ▶ Most assembler instructions represent machine instructions one-to-one

- ▶ Pseudoinstructions
 - Can be used in assembly language
 - Need to be translated into other instructions for translation into machine code
 - Example

```

move $t0, $t1           →    add $t0, $zero, $t1
blt  $t0, $t1, L        →    slt $at, $t0, $t1
                                bne $at, $zero, L
    
```

- `$at` (register 1): assembler temporary



- ▶ **Assembler (or compiler)**
 - Translates program into machine instructions
 - Provides information for building a complete program from the pieces

- ▶ **Object module**
 - **Header**
 - Described contents of object module
 - **Text segment**
 - Translated instructions (machine code)
 - **Static data segment**
 - Data allocated for the life of the program
 - **Relocation info**
 - For contents that depend on absolute location of loaded program
 - **Symbol table**
 - Remaining labels (e.g. external refs)
 - **Debug info**
 - For associating with source code



Linking Object Modules

- ▶ Programs consist of several procedures
- ▶ Procedures may be compiled and assembled independently
- ▶ These object modules need to be linked together
- ▶ Problem
 - Machine code needs to be relocated
 - External labels pointing to other procedures are not yet resolved
- ▶ That's the job of the linker or link editor (Binder)
- ▶ Linker produces an executable image
 - Merges segments
 - Resolve labels (determine their addresses)
 - Patch location-dependent and external labels/refs
- ▶ Could leave location dependencies for fixing by a relocating loader
- ▶ But with virtual memory, no need to do this
 - Program can be loaded into absolute location in virtual memory space



► Load from image file on disk into memory

1. Read header to determine segment sizes
2. Create virtual address space
3. Copy text and initialize data into memory
 - Or set page table entries so they can be faulted in
4. Set up arguments on stack
5. Initialize registers (including \$sp, \$fp, \$gp)
6. Jump to startup routine
 - Copy arguments to \$a0, ... and call main
 - When main returns, do syscall “exit”



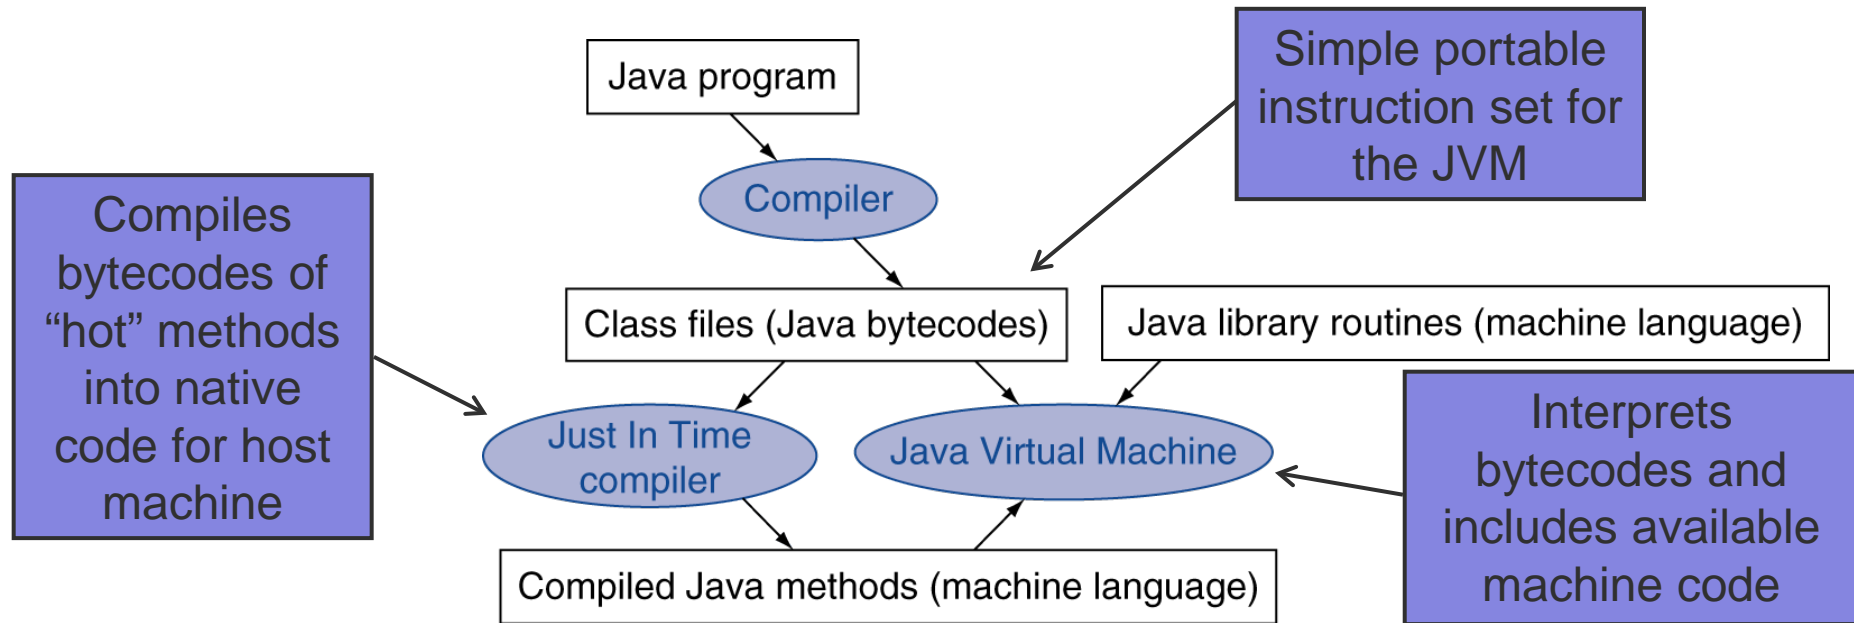
- ▶ Only link/load library procedure when it is called
 - Requires procedure code to be relocatable
 - Avoids image bloat caused by static linking of all (transitively) referenced libraries
 - Automatically picks up new library versions
 - Object module of such a library: dynamically linked library (DLL)

- ▶ First DLLs: all procedures of a library linked, requiring again large memory space

- ▶ Lazy linkage: links procedures at first call
 - More effort at first call
 - Less memory space needed



Starting Java Applications



- ▶ Interpreter = program simulating an ISA
- ▶ JVM interpretes Java bytecode
- ▶ Much slower than machine code

- ▶ JIT compiles Java bytecode into machine code during interpretation and saves it for future execution
- ▶ JVM also calls machine code from libraries and machine code bytecode compiled by JIT







- ▶ Instruction set architecture
- ▶ The MIPS ISA
- ▶ Calling procedures in MIPS
- ▶ Dealing with character sets
- ▶ Translating and starting a program
- ▶ Effect of compiler optimization
- ▶ Other ISAs
- ▶ Concluding remarks



Effect of Compiler Optimization

Compiled with gcc for Pentium 4 under Linux






- The unoptimized code has the best CPI, the O1 version has the lowest instruction count, but the O3 version is the fastest. Why?



Effect of Language and Algorithm

- ▶ Sort 100,000 words
 - with the array initialized to random values
- ▶ on a Pentium 4
 - with a 3.06 GHz clock rate,
 - a 533 MHz system bus,
 - with 2 GB of DDR SDRAM,
 - using Linux version 2.4.20





- ▶ Instruction count and CPI are not good performance indicators in isolation
- ▶ Compiler optimizations are sensitive to the algorithm
- ▶ Java/JIT compiled code is significantly faster than JVM interpreted
 - Comparable to optimized C in some cases
- ▶ Nothing can fix a dumb algorithm!



- ▶ Replaces a function call with the body of the function to be called (the callee)
- ▶ A compiler may do that before translation to assembler code
- ▶ Saves the effort of
 - Passing arguments and return values
 - Jumping
 - Saving and restoring registers
- ▶ May improve time usage at runtime
- ▶ May possibly increase the final size of the binary executable
 - Larger source code
- ▶ Widely used optimization technique in compilers



- ▶ `int b; c=b+5`
 - `b` is a variable for an integer
 - `&b` is the memory address of that variable
- ▶ `int *x; x=&b; c=*x+5;`
 - `x` is a variable for a address of a memory location holding an integer (**pointer**)
 - `*x` returns the value of that location (**dereferencing**)
- ▶ Pointer arithmetic
 - „sizeof(int)“ is size of an integer in bytes
 - `x++`: pointer `x` is increased by `sizeof(int)`, not the value of `*x`!
- ▶ Connection with arrays: `int a[]`
 - Several items of same type consecutively stored in memory
 - `int *x=a`; base address array
 - Value and address of `i`-th position in array
 - `a[i]`, `&a[i]`; `*(x+i)`, `(x+i)`

...	...
<u>a[4]=4</u>	a+16
<u>a[3]=5</u>	a+12
<u>a[2]=7</u>	a+8
<u>a[1]=2</u>	a+4
<u>a[0]=3</u>	a+0



- ▶ Indexing general arrays (e.g. of structs) involves
 - Multiplying index by element size
 - Adding to array base address
- ▶ Pointers correspond directly to memory addresses
 - Can avoid indexing complexity
- ▶ More efficient code possible with pointers (see next slide)
- ▶ But compiler optimization
 - “strength reduction” effects usage of shift instead of multiply by 2^n
 - “elimination of induction variables” converts usage of array indexing to usage of pointers
- ▶ Recommendation
 - Write more readable code using array indexing
 - Let the compiler do the tricky stuff



Example: Clearing an Array

- ▶ Arrays and pointers lead to almost the same assembly code
- ▶ But blue instructions are required to be in the loop only for arrays

```
clear1(int array[], int size) {
    int i;
    for (i = 0; i < size; i += 1)
        array[i] = 0;
}
```

```

    move $t0,$zero    # i = 0
loop1: sll  $t1,$t0,2   # $t1 = i * 4
    add  $t2,$a0,$t1   # $t2 =
                        # &array[i]
    sw   $zero, 0($t2) # array[i] = 0
    addi $t0,$t0,1     # i = i + 1
    slt  $t3,$t0,$a1    # $t3 =
                        # (i < size)
    bne  $t3,$zero,loop1 # if (...)
                        # goto loop1
```

```
clear2(int *array, int size) {
    int *p;
    for (p = &array[0]; p < &array[size];
        p = p + 1)
        *p = 0;
}
```

```

    move $t0,$a0      # p = & array[0]
    sll  $t1,$a1,2     # $t1 = size * 4
    add  $t2,$a0,$t1   # $t2 =
                        # &array[size]
loop2: sw   $zero,0($t0) # Memory[p] = 0
    addi $t0,$t0,4     # p = p + 4
    slt  $t3,$t0,$t2   # $t3 =
                        # (p<&array[size])
    bne  $t3,$zero,loop2 # if (...)
                        # goto loop2
```

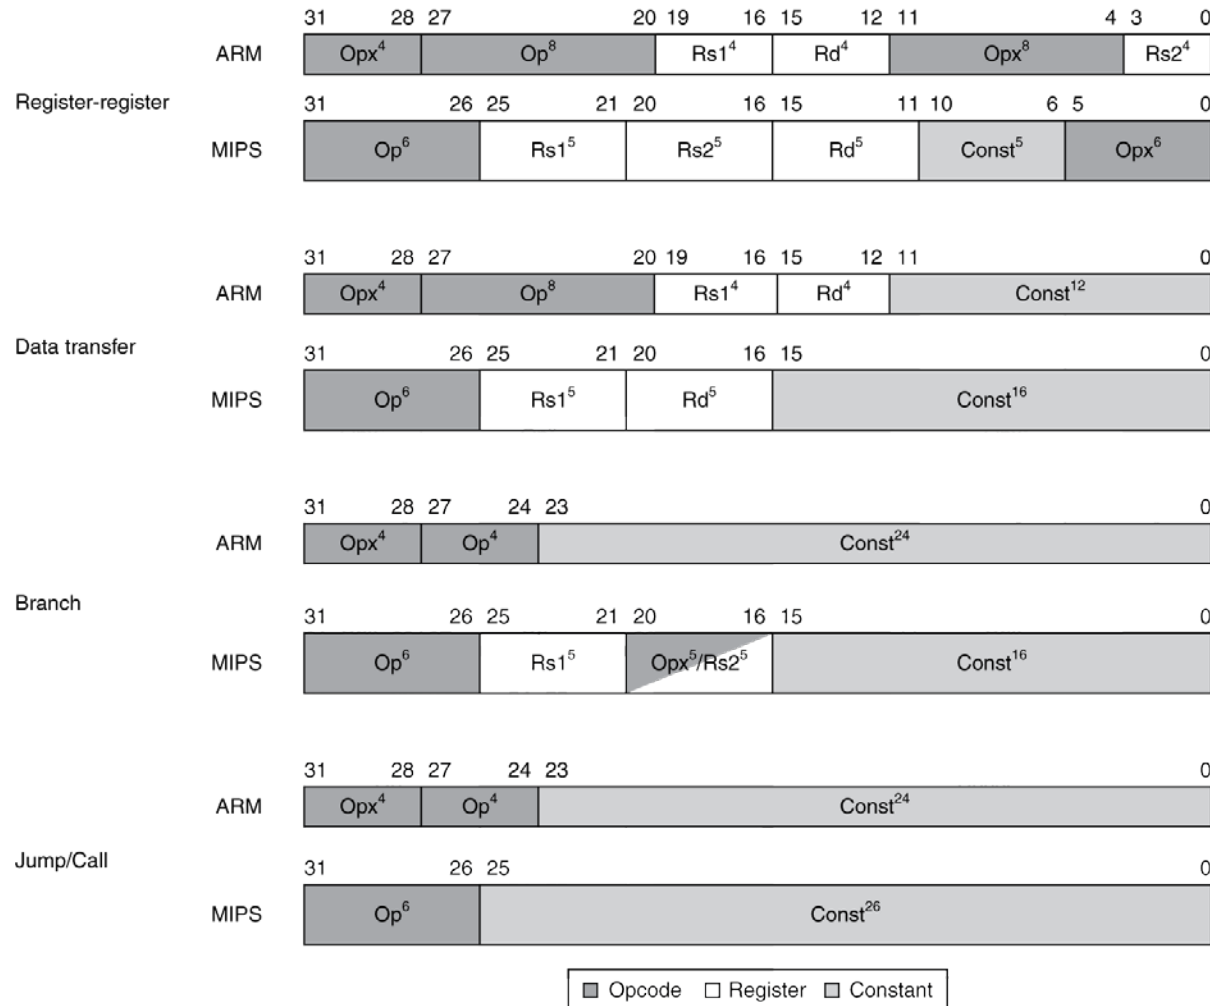


- ▶ Instruction set architecture
- ▶ The MIPS ISA
- ▶ Calling procedures in MIPS
- ▶ Dealing with character sets
- ▶ Translating and starting a program
- ▶ Effect of compiler optimization
- ▶ Other ISAs
- ▶ Concluding remarks



- ▶ ARM (Advanced RISC Machines)
- ▶ The most popular embedded core, widely used in smartphones
- ▶ Basic set of instructions similar to MIPS

	ARM	MIPS
Date announced	1985	1985
Instruction size	32 bits	32 bits
Address space	32-bit flat	32-bit flat
Data alignment	Aligned	Aligned
Data addressing modes	9	3
Registers	15 × 32-bit	31 × 32-bit
Input/output	Memory mapped	Memory mapped



► ARM uses different principle for branch instructions



- ▶ IA-32: Intel Architecture 32 bit (32 bit x86 architecture)
 - Ex. 80386 / 386 / i386
- ▶ General purpose registers (GPRs) from 80386 on
- ▶ Words
 - 16 bit = word, 32 bit = double word, 64 bit = quad word
- ▶ Instructions
 - Significantly more than in MIPS
 - More complex
 - Varying length instruction format (1 – 15 bytes)
- ▶ Complex instruction set makes implementation difficult
 - Hardware translates instructions to simpler microoperations
 - Simple instructions: 1–1
 - Complex instructions: 1–many
 - Microengine similar to RISC
 - Market share makes this economically viable
- ▶ Comparable performance to RISC
 - Compilers avoid complex instructions
- ▶ Evolution with backward compatibility
 - If Intel didn't extend with compatibility, its competitors would!
 - Technical elegance ≠ market success



- ▶ 8080 (1974): 8-bit microprocessor
 - Accumulator, plus 3 index-register pairs
- ▶ 8086 (1978): 16-bit extension to 8080
 - Complex instruction set (CISC)
- ▶ 8087 (1980): floating-point coprocessor
 - Adds FP instructions and register stack
- ▶ 80286 (1982): 24-bit addresses, MMU
 - Segmented memory mapping and protection
- ▶ 80386 (1985): 32-bit extension (now IA-32)
 - Additional addressing modes and operations
 - Paged memory mapping as well as segments
- ▶ i486 (1989): pipelined, on-chip caches and FPU
 - Compatible competitors: AMD, Cyrix, ...
- ▶ Pentium (1993): superscalar, 64-bit datapath
 - Later versions added MMX (Multi-Media eXtension) instructions
 - The infamous FDIV bug




- ▶ Pentium Pro (1995), Pentium II (1997)
 - New microarchitecture (see Colwell, *The Pentium Chronicles*)
- ▶ Pentium III (1999)
 - Added SSE (Streaming SIMD Extensions) and associated registers
- ▶ Pentium 4 (2001)
 - New microarchitecture
 - Added SSE2 instructions
- ▶ AMD64 (2003): extended architecture to 64 bits
- ▶ EM64T – Extended Memory 64 Technology (2004)
 - AMD64 adopted by Intel (with refinements)
 - Added SSE3 instructions
- ▶ Intel Core (2006)
 - Added SSE4 instructions, virtual machine support
- ▶ AMD64 (announced 2007): SSE5 instructions
 - Intel declined to follow, instead...
- ▶ Advanced Vector Extension (announced 2008)
 - Longer SSE registers, more instructions
- ▶ ...



- ▶ Instruction set architecture
- ▶ The MIPS ISA
- ▶ Calling procedures in MIPS
- ▶ Dealing with character sets
- ▶ Translating and starting a program
- ▶ Effect of compiler optimization
- ▶ Other ISAs
- ▶ Concluding remarks



- ▶ Powerful instruction \Rightarrow higher performance
 - Fewer instructions required
 - But complex instructions are hard to implement
 - May slow down all instructions, including simple ones
 - Compilers are good at making fast code from simple instructions
- ▶ Use assembly code for high performance
 - But modern compilers are better at dealing with modern processors
 - More lines of code \Rightarrow more errors and less productivity
- ▶ Backward compatibility
 - \Rightarrow instruction set doesn't change
 - But they do accrete more instructions





- ▶ Sequential words are not at sequential addresses
 - Increment by 4, not by 1!

- ▶ Pointers to local variables (e.g. in procedures)
 - Become invalid when stack popped
 - But address can be passed back via an argument
 - Don't do that!



► Design principles

1. Simplicity favors regularity \Rightarrow fixed instruction formats
2. Smaller is faster \Rightarrow only a few registers
3. Make the common case fast \Rightarrow PC-relative addressing
4. Good design demands good compromises \Rightarrow immediate addressing for small constants help to keep fixed length instruction formats

► Layers of software/hardware

- Compiler, assembler, hardware

► MIPS: typical of RISC ISAs

► X86: not typical RISC ISA (more complex)